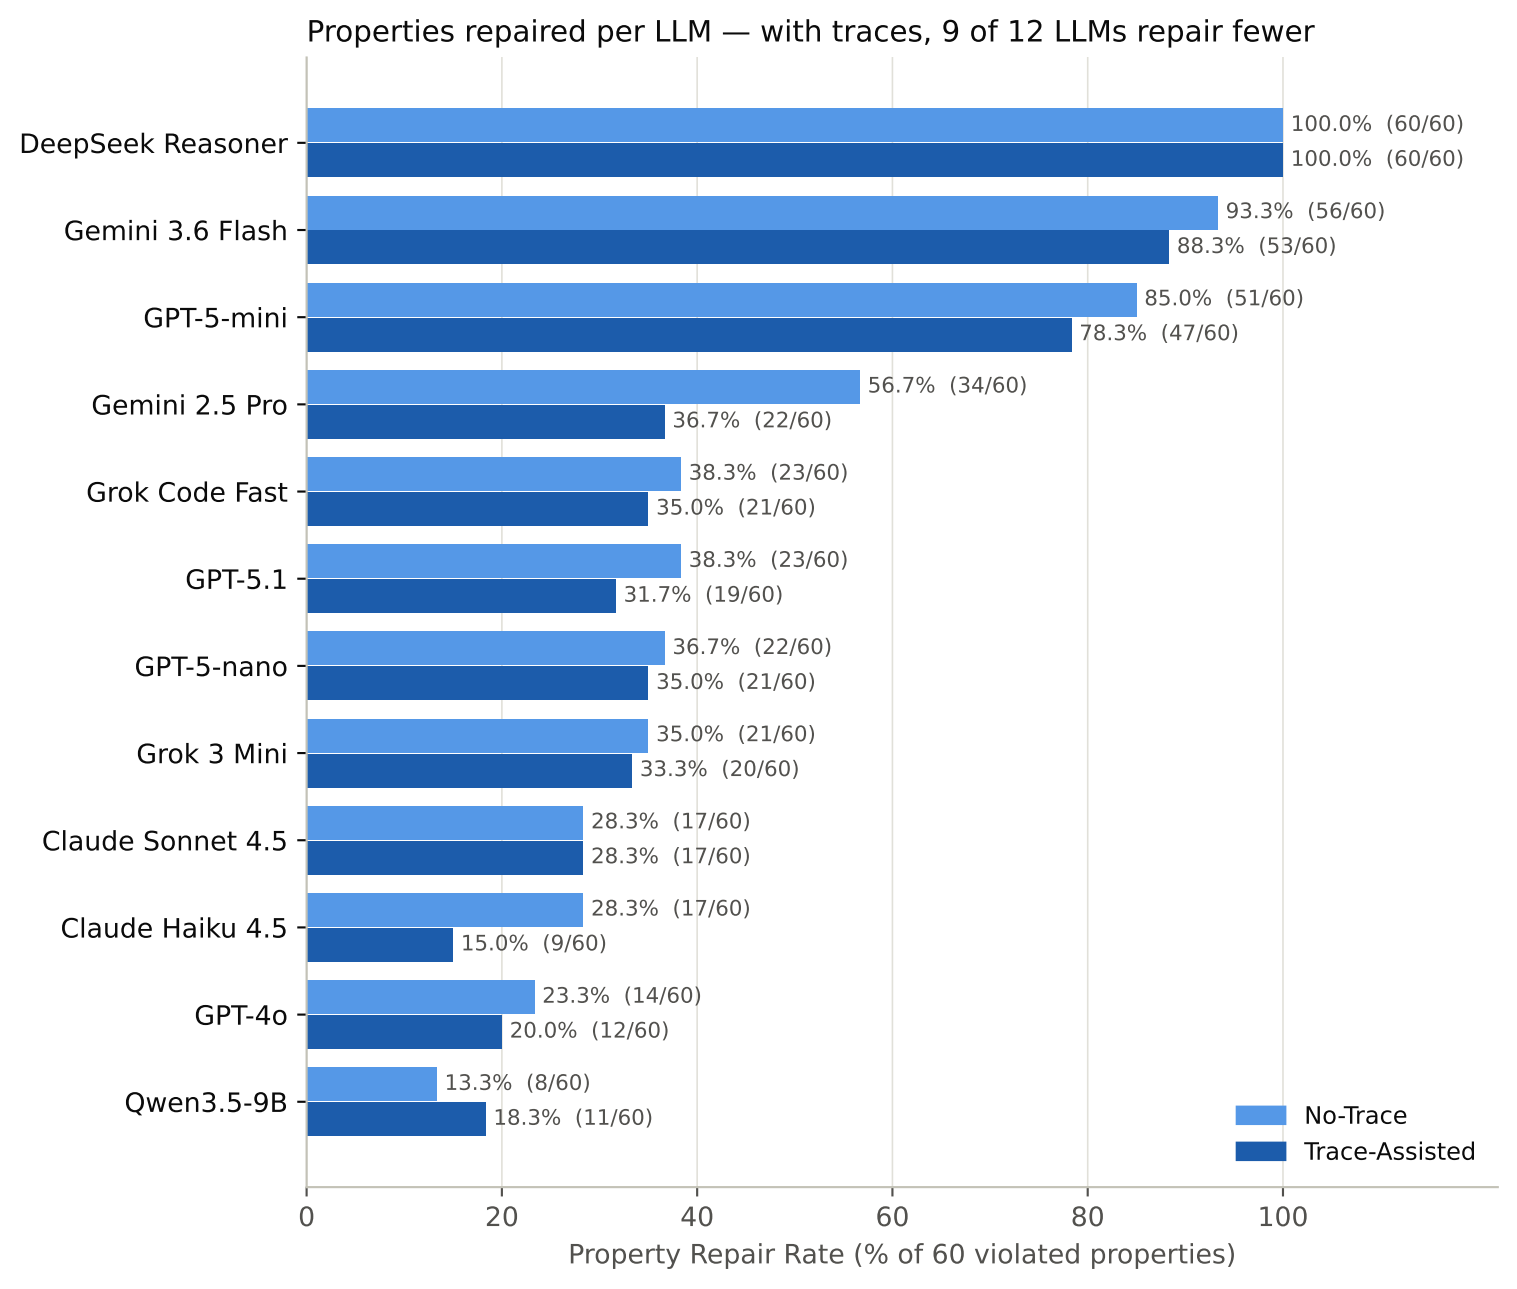
Properties repaired per LLM — with traces, 9 of 12 LLMs repair fewer
DeepSeek Reasoner
100.0% (60/60)
100.0% (60/60)
Gemini 3.6 Flash
93.3% (56/60)
88.3% (53/60)
GPT-5-mini
85.0% (51/60)
78.3% (47/60)
Gemini 2.5 Pro
56.7% (34/60)
36.7% (22/60)
Grok Code Fast
38.3% (23/60)
35.0% (21/60)
GPT-5.1
38.3% (23/60)
31.7% (19/60)
GPT-5-nano
36.7% (22/60)
35.0% (21/60)
Grok 3 Mini
35.0% (21/60)
33.3% (20/60)
Claude Sonnet 4.5
28.3% (17/60)
28.3% (17/60)
Claude Haiku 4.5
28.3% (17/60)
15.0% (9/60)
GPT-4o
23.3% (14/60)
20.0% (12/60)
Qwen3.5-9B
13.3% (8/60)
18.3% (11/60)
No-Trace
Trace-Assisted
0
20
40
60
80
100
Property Repair Rate (% of 60 violated properties)